

# CRScore: Grounding the Evaluation of Code Review Comments in Code Claims and Smells

Atharva Naik   Marcus Alenius   Daniel Fried   Carolyn Rosé

Language Technologies Institute

Carnegie Mellon University

{arnaik, malenius, dfried, cprose}@cs.cmu.edu

## Abstract

The task of automated code review has recently gained a lot of attention from the machine learning community. However, current review comment evaluation metrics rely on comparisons with a human-written reference for a given code change (also called a *diff*), even though code review is a one-to-many problem like generation and summarization with many “valid reviews” for a diff. To tackle these issues we develop a CRScore — a reference-free metric to measure dimensions of review quality like conciseness, comprehensiveness, and relevance. We design CRScore to evaluate reviews in a way that is grounded in claims and potential issues detected in the code by LLMs and static analyzers. We demonstrate that CRScore can produce valid, fine-grained scores of review quality that have the greatest alignment with human judgment (0.54 Spearman correlation) and are more sensitive than reference-based metrics. We also release a corpus of 2.6k human-annotated review quality scores for machine-generated and GitHub review comments to support the development of automated metrics.

## 1 Introduction

Code Review is an essential peer analysis quality control tool for software engineers to ensure that source code is free of bugs and upholds standards (McIntosh et al., 2014; Bavota and Russo, 2015). Software engineers prefer lightweight, async reviews like GitHub’s review comment feature over formal in-person reviews (Beller et al., 2014; Badampudi et al., 2023) leading to the creation of benchmarks for automated generation of natural language (NL) review comments (Li et al., 2022; Tufano et al., 2022). However, these benchmarks use reference-based evaluation metrics like BLEU (Papineni et al., 2002), which have been shown to have low validity (Reiter, 2018; Evtikhiev et al., 2023), especially when paired with limited and low-quality references.

Code review fundamentally is a one-to-many problem, where a given diff can have multiple possible issues that can be tackled by a review. Having a limited number of reference reviews (e.g. one per diff in CodeReviewer (Li et al., 2022)) leads to unfairly low scores with reference-based metrics. For example for the diff shown in Figure 1, the ground truth review focuses on whether the `ToHexString()` function could be a performance issue. However, the model-generated review focuses on the `ToHexString().Equals("0000000000000000")` being an odd condition with a scenario where it’s triggered being unlikely, which is also a valid review for the diff. However, the BLEU score value for the model-generated review is very low at 0.0458 due to poor n-gram overlap. Additionally, the references can also be low-quality with missing context, as shown in Table 4. Such low-quality references paired with reference-based metrics can harshly and unfairly penalize models.

Motivated by these drawbacks we propose CRScore an automated but reference-free evaluation metric that uses dimensions of review quality from prior work (Piorkowski et al., 2020; Turzo and Bosu, 2024). Namely, *Comprehensiveness* – does the review convey all the necessary information? *Conciseness* – does the review only convey the necessary information in an efficient way and *Relevance* – is all the information on topic. We operationalize our metric through a two-step process: ① generate a list of *pseudo-references* spanning information like possible claims, issues, and implications of a code change, ② use semantic textual similarity (STS) to align parts of the review to the pseudo-references. To generate the pseudo-references we use a neuro-symbolic approach that combines **Large Language Models** (LLMs) and **Code Analysis Tools** (CATs) that can detect formatting errors, faulty design patterns (code smells Rasheed et al. 2024), etc. We combine these meth-

ods for more exhaustive pseudo-references and to overcome drawbacks of each method (section A.1).

Finally, we demonstrate the validity of CRScore by human evaluation of the quality of pseudo-references and measuring the alignment of our metrics with human judgment. We show that a large number (82.6%) of pseudo-references generated by CRScore are correct and it has the greatest alignment (Spearman correlation 0.5431) and sensitivity (Figure 8) to the human judgment of review quality, compared to the reference-based metrics.

**Contributions.** (1) We propose an automated reference-free metric that combines the advantages of LLMs and CATs to measure review quality along fine-grained dimensions. (2) We collect human annotations of pseudo-reference quality and review conciseness, comprehensiveness, and relevance to validate our approach. We plan to make them publicly available as a resource for the development of automated metrics. (3) We benchmark several LLMs of code for code review on the CodeReviewer dataset using CRScore and compare it with reference based metrics.

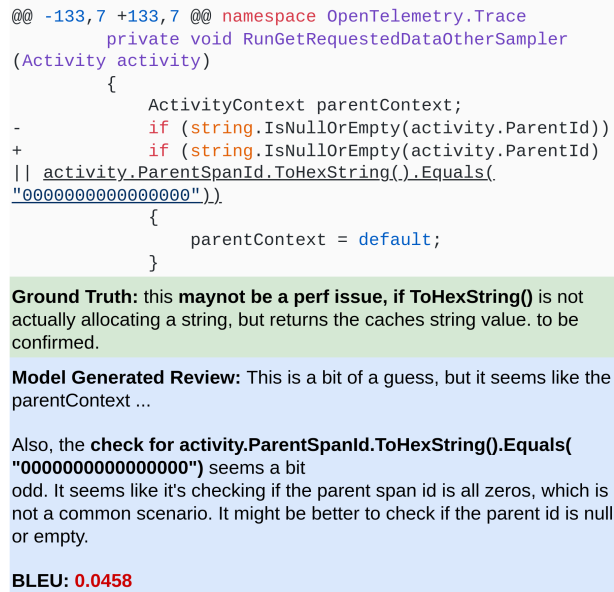


Figure 1: Example diff with multiple valid reviews. The ground truth and model-generated reviews focus on different topics like the performance of the added check and how likely it is to be triggered. However, a reference-based metric like the BLEU score assigns this review a low score of 0.0458.

## 2 Related Work

In this section, we summarize the limitations of reference-based evaluation, the need for better

code review evaluation metrics, inspiration from reference-free evaluation for other tasks, and how code smells can be leveraged to evaluate reviews.

### 2.1 Reference Based Evaluation Metrics:

Reference-based metrics like BLEU (Papineni et al., 2002), ROUGE (Ganesan, 2018), and BERTScore (Zhang et al., 2020) have seen widespread adoption for several text generation tasks like machine translation and summarization due to their convenience. While BLEU and ROUGE and similar metrics like character F-score (Popović, 2015, 2017) use n-gram overlap between the reference and candidate text, metrics like BERTScore (Zhang et al., 2020) try to capture the semantic similarity. However prior studies have shown metrics like BLEU to have low validity (overlap with human judgment) and reliability (Reiter, 2018; Evtikhiev et al., 2023). Meanwhile, BERTScore can fall for candidates with errors that are lexically and stylistically similar to references (Hanna and Bojar, 2021).

### 2.2 Code Review Evaluation

Due to the high time and resource demands of code review automated approaches have gained popularity (Yang et al., 2024). Tufano et al. (2021, 2022); Li et al. (2022) proposed the most popular datasets for code review tasks like detecting code changes that require review, generating review comments, and refactoring code. However, both of these datasets used reference-based metrics and thus suffer from the issues highlighted in the previous section. While many studies have focused on modeling methods for code review tasks (Pornprasit and Tantithamthavorn, 2024; Lu et al., 2023; Dong-Kyu, 2024; Fan et al., 2024; Yu et al., 2024; Lin et al., 2024), these studies either retain the same reference-based automated metrics like BLEU score (Papineni et al., 2002) for evaluating review comments or use human evaluation. In this work, we show that the use of reference-based metrics combined with noisy references fails to capture human notions of review quality. Hence we propose CRScore the first reference-free metric for code review comment generation to overcome the limitations of these metrics.

### 2.3 Reference Free Evaluation

Reference-free evaluation metrics have been proposed for various generative tasks like dialog to deal with the issue of multiple valid outputs. These

metrics broadly target the evaluation of various “quality dimensions” like relevance, informativeness, etc. VIFIDEL (Madhyastha et al., 2019) InfoMetIC (Hu et al., 2023), ClipScore (Hessel et al., 2021) evaluate dimensions like faithfulness, informativeness, and relevance, respectively for image captioning. FED (Mehri and Eskenazi, 2020a), and USR (Mehri and Eskenazi, 2020b) have been proposed for evaluating dialog across dimensions like informativeness, relevance, and overall quality. Finally, studies on the helpfulness of software documentation like (Piorkowski et al., 2020) report the usefulness of quality dimensions like conciseness, completeness, relevance, and supporting evidence while studies on code review like (Turzo and Bosu, 2024) identify concise and understandable reviews to be helpful. The common trend across these studies is some notion of conciseness, informativeness (comprehensiveness in our work), and relevance being useful, prompting us to focus on them. Additionally, our metric is also similar to InfoMetIC which has text precision, visual recall, and overall quality. However, none of these studies target code reviews, making our work the first automated reference-free metric for code reviews.

## 2.4 Code Smell Detection

“Code smells” (Fowler, 1997) are symptoms of design flaws and bad practices (also called *anti-patterns*) that can snowball into maintainability issues. Detecting code smells automatically has been traditionally done by analysis-based heuristic approaches (Tsantalis et al., 2008; Paiva et al., 2017; Liu and Zhang, 2017), with machine learning (Sandouka and Aljamaan, 2023) and transfer learning (Sharma et al., 2021) based approaches being proposed for learning more complex heuristics. Recent approaches have leveraged LLMs via prompting (Liu et al., 2024) and agents Rasheed et al. (2024) to achieve further improvement or tackle repository-level code smell detection. However, these approaches are limited in the types of smells they target (Liu et al., 2024; Sandouka and Aljamaan, 2023), training data requirements (Zhang et al., 2024), or lack comprehensive evaluation (Rasheed et al., 2024). Also, code smells differ across programming languages (Abidi et al., 2019), and transfer learning approaches can only be leveraged for similar languages (Sharma et al., 2021). Due to these limitations of learning-based methods and to mitigate the self-selection bias of LLMs

(section A.1) we use code analysis tools.

## 3 Operationalizing CRScore

Motivated by the one-to-many nature of code reviews, noisy references, and the pitfalls of reference-based automated metrics we develop CRScore — a reference-free, quality dimension-based automated metric. As shown in Figure 2, instead of relying on explicit references, our metric generates “pseudo-references” from the code change spanning claims, implications, and issues or smells that could hurt maintainability down the line — in other words some topics that a review should address (Rasheed et al., 2024). Then we use semantic textual similarity (STS) measures to quantify how much these topics are addressed by a code review as shown in Figure 2 through the lens of three quality dimensions: conciseness, comprehensiveness, and relevance. They capture review quality in a way similar to precision, recall, and f-score for classification and retrieval. We describe the three main components of our framework — the quality dimensions, pseudo-reference generation, and similarity measurement — below.

### 3.1 Quality Dimensions

We pick a subset of the quality dimensions proposed by Piorkowski et al. (2020) for software documentation evaluation and adapt them to code review. We pick dimensions of “completeness” and “conciseness” Piorkowski et al. (2020) to capture and strike a balance between comprehensive reviews with a lot of detail and concise, minimalist reviews. To ensure that reviews are comprehensive while minimizing irrelevant information we propose a precision-like measure called *conciseness*, a recall-like measure *comprehensiveness*, and a combined measure of the overall quality or *relevance* similar to the f-measure (and “relevance” dimension from Piorkowski et al. 2020). Additionally, we also do a human evaluation of the validity of the pseudo-references focusing on aspects like supporting evidence (Piorkowski et al., 2020).

### 3.2 Pseudo Reference Generation

To generate pseudo references we develop an LLM-based pipeline for generating **claims** about the code changes on two levels of abstraction: 1) Low-level changes and 2) High-level “implications” of the change. To allow for reproducibility and ease of deployment we use a 6.7B parameter open source

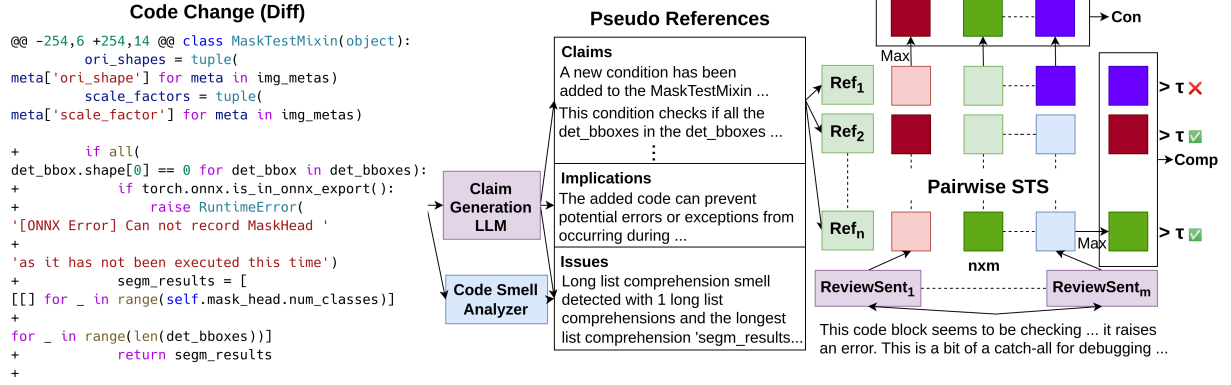


Figure 2: **Operationalization of CRScore:** Our metric first generates pseudo-references for the diff — claims, implications and issues. Then each pseudo-reference is embedded by a sentence transformer along with each review sentence and the pairwise semantic textual similarity (STS) is computed. The high similarity threshold  $\tau$  is used to compute the Con and Comp metrics whose harmonic mean yields the Rel score.

LLM, Magicoder (Wei et al., 2023). We fine-tune Magicoder-S-DS-6.7B for this task using task-specific data produced by a stronger model (GPT-4) as shown in Figure 3. We generate claims by prompting GPT-4 for a random subset of 1k code changes from the CodeReviewer validation data.

In addition to the claims we also utilize program analysis tools that can detect **issues** and “code smells” (characteristics that indicate deeper flaws in programs as shown in Table 13) like PySmell<sup>1</sup> as well as static analyzers like PMD<sup>2</sup> and JSHint<sup>3</sup>. They can detect unused variables, unnecessary object creation, syntax errors, leaking variables, type conversion issues, etc. These tools target Python, Java, and Javascript respectively, and can use rule-based analysis to detect issues at a file or repository level. They complement the aspects that might be missed by an LLM-based analysis. We combine the code smells and claims to get the final set of pseudo references.

### 3.3 Computing Similarity with Pseudo References:

We use a Sentence Transformer (Reimers, 2019) model (mx-bai-embed-large-v1<sup>4</sup>) to compute Semantic Textual Similarity (STS) between pseudo-references and review sentences. The pairwise similarities are then used to compute the conciseness, comprehensiveness, and relevance as shown by

equation 1, 2, and 3. We picked this model because it has the best performance as of July 2024 on English STS on the MTEB benchmark (Muenighoff et al., 2022) for models with less than 1B parameters.

We start by computing token embeddings for the pseudo-references (p-refs) and review sentences (r-sents), excluding the stopwords and pooling the rest of the token embeddings to build representations for the whole sentence. Then we compute STS scores with these sentence embeddings via pairwise cosine similarity  $s()$  between the p-refs ( $\mathcal{P}$ ) and r-sents ( $\mathcal{R}$ ). The Conciseness ( $Con$ ) which is computed as:

$$Con = \frac{\sum_{r \in \mathcal{R}} I[\max_{p \in \mathcal{P}} s(c, r) > \tau]}{|\mathcal{R}|} \quad (1)$$

represents the fraction of r-sents from the model-generated review with greater similarity to any p-ref above a threshold  $\tau$ . Here,  $I$  is an indicator variable such that:

$$I[x] = \begin{cases} 1, & \text{if } x \text{ is true} \\ 0, & \text{otherwise} \end{cases}$$

$Con$  resembles precision as it captures the fraction of r-sents (candidate set) that are “on topic” concerning the p-refs (reference/gold set). The Comprehensiveness ( $Comp$ ) computed as:

$$Comp = \frac{\sum_{p \in \mathcal{P}} I[\max_{r \in \mathcal{R}} s(c, r) > \tau]}{|\mathcal{P}|} \quad (2)$$

represents the fraction of p-refs that have greater similarity to any of the r-sents than the threshold  $\tau$ .

<sup>1</sup><https://github.com/whyjay17/Pyscent>

<sup>2</sup><https://pmd.github.io/>

<sup>3</sup><https://jshint.com/>

<sup>4</sup><https://huggingface.co/mixedbread-ai/mx-bai-embed-large-v1>



This metric resembles recall as it captures the fraction of the p-refs (gold set) covered by the model-generated r-sents (candidate set). There is a trade-off between conciseness and comprehensiveness just like precision and recall, and to capture the trade-off between these metrics like the F1 score, we define the Overall Relevance  $Rel$  as the harmonic mean of  $Con$  and  $Comp$ :

$$Rel = \frac{2 \cdot Con \cdot Comp}{Con + Comp} \quad (3)$$

Based on the definition of  $I[x]$  all of our metrics range in value from  $[0, 1]$ .

**Selecting threshold  $\tau$  for high similarity:** For our similarity threshold  $\tau$  we use the average similarity between a pseudo-reference and the review sentence talking about it. To compute it we use the distribution of pseudo-reference and review similarity scores for the best model (GPT-3.5 in our case). Starting with the CodeReviewer test set we exclude the data used for collecting human annotations. For the remaining 9869 instances, we compute the similarity between each review sentence and pseudo-reference for its corresponding code change. Then we associate each review sentence with the most similar pseudo-reference to capture correspondence. Finally, we average the similarity score value across all review sentences, giving us a value for the threshold<sup>5</sup>  $\tau_{best} = 0.7314$ .

## 4 Validating CRScore

To be a valid metric, CRScore needs to satisfy a few properties. Firstly the generated pseudo-references should have few errors and unverifiable claims and be as exhaustive as possible. Additionally, we want our metric dimension scores ( $Con$ ,  $Comp$ , and  $Rel$ ), especially  $Rel$  to correlate with human judgment for the corresponding dimension. However, arguably most importantly we want our metric to rank a diverse set of review generation systems similar to how humans would rank them based on overall review relevance. In the subsequent sections, we describe how we design the experiments to collect these annotations (section 4.1, 4.2), choose a set of systems to be rated by humans (section 4.3), and set up reference-based metrics (section 4.4) for comparison against CRScore.

<sup>5</sup>We also show that computing the threshold with ground truth reviews yields a comparably good threshold  $\tau_{GT} = 0.6576$  with respect to correlation with human annotations.

### 4.1 Rating Quality of Pseudo-References

To show that LLM-generated pseudo-references used by the CRScore evaluation pipeline are high-quality, we gather annotations for their quality, capturing incorrect, unverifiable, and missing claims. To rate the quality of the pseudo-references, we had two trained human annotators judge their quality for 100 randomly sampled code changes each in Python, Java, and Javascript (300 total). The annotators were asked to code the pseudo-references as 1 (correct based on evidence), 0 (incorrect based on evidence), and -1 (unverifiable due to lack of evidence). They were also asked to add any pseudo-references about issues/claims not covered by the pseudo-references. We report the fraction of correct claims (accuracy), incorrect claims (error rate), unverifiable claims (unverifiable rate), and missing claims (missing rate). If  $N_c$ ,  $N_u$ ,  $N_i$ , and  $N_a$  represent the number of correct, unverifiable, incorrect, and added claims then each of these rates can be calculated as:

$$Accuracy = \frac{N_c}{N_c + N_u + N_i}$$

$$Error Rate = \frac{N_i}{N_c + N_u + N_i}$$

$$Unverifiable Rate = \frac{N_u}{N_c + N_u + N_i}$$

$$Missing Rate = \frac{N_a}{N_c + N_u + N_i}$$

Based on these expressions: Accuracy + Unverifiable Rate + Error Rate = 1. The results for each language are shown in table 1. The annotators follow guidelines laid out in a codebook (section D.1) which includes examples for each category. We measured the coding reliability of our approach by collecting annotations from both annotators on a common set of 100 pseudo-references. These annotations yielded a Cohen Kappa of 0.804 which indicates great inter-annotator reliability (Landis and Koch, 1977). Also, we only evaluate the LLM-generated claims here as we know the static code analysis tools are rule-based and reliably correct.

### 4.2 Rating Review Quality Dimensions

To show that CRScore aligns with the human judgment of review quality along the proposed dimensions: comprehensiveness, conciseness, and relevance we gather annotations from the same annotators on reviews generated by 9 systems (section 4.3)

and the ground truth references. We use the same 300 code changes used in section 4.1.

The raters again follow a codebook (section D.2) that contains guidelines and examples for annotating each dimension using the pseudo-references on a 5-point Likert scale with 5 being the highest and 1 the lowest. The raters use the updated set of pseudo-references based on the previous round of annotations. Incorrect and unverifiable references are removed and missing claims are added. We also add the pseudo-references (issues and code smells) generated by the code analysis tools for each code change. Raters are also asked to annotate any claims they find unnecessary for a given code change which are then excluded while rating the review on each quality dimension. Again we measure the reliability of the codebook by collecting annotations from the two raters, this time on a common set of 100 reviews. We compute the Krippendorff’s alpha reliability coefficient (Krippendorff, 2018) for each dimension, yielding values of 0.8868, 0.8505 and 0.8806 for conciseness, comprehensiveness, and relevance. Alpha values > 0.8 are generally considered reliable.

We measure the agreement between CRScore (Rel) and the reference-based metrics with the human-annotated relevance Likert scores using Kendal and Spearman Rank correlation (Table 2). We also compute the correlation between the system rankings generated by the metric and human annotations (Table 3).

### 4.3 Review Generation Systems

To test the ability of CRScore to rank code review systems of varying capabilities, we choose a diverse set of review generation systems. They span various parameter sizes, pre-training, and domain-specific fine-tuning strategies for code review:

**Simple baselines:** We create two simple baselines a BM-25 retriever and an LSTM as described in section D.3. We choose these models with the expectation that they will likely perform the worst, to see if our metric assigns them a low score.

**CodeReviewer:** We pick the CodeReviewer model from (Li et al., 2022) as it is a transformer-based model trained on code review-specific data and objectives.

**Open source LLMs:** We prompt several open-source LLMs in a few-shot manner with a fixed set of three example code changes and review pairs from the validation set. We use LLMs in the 3-13B

parameter range: Stable-Code-3B (Pinnaparaju et al.), DeepSeekCoder-6.7B (Guo et al., 2024), Magicoder-6.7B (Wei et al., 2023), CodeLLaMA-7B and 13B (Roziere et al., 2023) and LLaMA-3-8B (AI@Meta, 2024).

**Closed source LLMs:** We prompt closed-source LLMs like GPT-3.5 in a manner similar to the open-source LLMs.

### 4.4 Reference-based Metrics

We pick commonly used reference-based metrics for code review and other text generation tasks to compare with CRScore:

**BLEU** (Papineni et al., 2002) measures the n-gram precision between the generated text and references with an additional brevity penalty to discourage short outputs. It is used for evaluation in both (Tufano et al., 2021) and CodeReviewer (Li et al., 2022). We report the results with and without stop word removal.

**Normalized Edit Distance** is a normalized Levenshtein distance used in prior work (Tufano et al., 2021; Baire et al., 2024) to measure the number of edits required to match candidate and target reviews or code.

**ROUGE-L F-measure** is a popular recall-oriented metric originally proposed for summarization and machine translation. We use the longest common subsequence-based sentence level f-measure implementation.<sup>6</sup>

**chrF:** (Popović, 2015) Is a machine translation metric which is essentially a character level F score computed using character level n-grams.

**chrF++:** (Popović, 2017) Is a variant of chrF that incorporates word n-grams as well and was found to correlate more with human assessment.

**BERTScore:** (Zhang et al., 2020) We use the BERTScore F1 measure to capture the semantic similarity between the reference review and the generated review.

## 5 Results

### 5.1 Validity of Pseudo-References

We show the rates of correct, incorrect, unverifiable, and missing claims as explained in section 4.1 in Table 1. The pseudo-references produced by our pipeline were relatively accurate according to the annotators with roughly 82.6% accuracy across the languages. The best performance is for

<sup>6</sup><https://pypi.org/project/rouge-score/>

Language	Accuracy	Error Rate	Unverifiable Rate	Missing Rate
Python	83.65	12.02	4.33	5.77
Java	79.09	13.22	7.69	8.89
Javascript	85.21	6.52	8.27	2.51

Table 1: Quality of pseudo-references, reported as the fraction of correct claims (accuracy), incorrect claims (error rate), unverifiable claims (unverifiable rate), and added claims (missing rate) based on human annotations.

Javascript and the worst performance is for Java. The most frequent issues in the code claims were incorrect claims (for Java and Python). Most of the errors were in code comprehension (“misreading the code”) and over-generalization (incorrect assumptions/generalizations made from the limited context, contradicting file level context). Some examples are shown in Table 11.

For Javascript, we observed unverifiable claims to be the biggest issue. There were a variety of claims about code efficiency, functionality, or performance made without evidence. E.g.: “However, this change could also potentially enable less strict or less strict-like behavior, depending on the context in which the function is used. This could make the code less efficient or less performant”.

Metric	$\tau$	$r_s$
BLEU	0.001	-0.0001
BLEU (without stopwords)	0.0425	0.0542
BERTScore	0.081	0.1083
Normalized Edit Distance	0.0193	0.0249
ROUGE-L	0.0757	0.0989
F-measure		
chrF	0.1484	0.1966
chrF++	0.1555	0.2057
<b>Rel (<math>\tau_{best}</math>) (Ours)</b>	<b>0.4567</b>	<b>0.5431</b>

Table 2: Comparing Kendal-Tau ( $\tau$ ) and Spearman Rank ( $r_s$ ) correlation of reference-based evaluation metrics and our reference-free relevance score (Rel) with human annotations for the relevance dimension. Results that don’t achieve statistical significance are grayed out.

Metric	$r_s$	$p$	$\tau$	$p$
BLEU	-0.3	0.433	-0.1667	0.612
BLEU (without stop)	-0.15	0.7	-0.0556	0.919
BERTScore	0.35	0.356	0.2222	0.477
Normalized Edit Distance	0.1667	0.668	0.0556	0.919
ROUGE-L	0.0167	0.966	0.556	0.919
F-measure				
chrF	0.4833	0.187	0.3889	0.18
chrF++	0.6	0.088	0.4444	0.119
<b>Rel (<math>\tau_{best}</math>) ours</b>	<b>0.95</b>	<b>8.7e-5</b>	<b>0.8889</b>	<b>2.4e-4</b>

Table 3: Spearman and Kendall Rank Correlations between system rankings produced by each metric and human annotations (Rel). Only our metric achieves a statistically significant and high correlation. Only LLaMA-3 is ranked lower than expected by our metric.

## 5.2 Validity of Review Quality Dimension Scores (Con, Comp and Rel)

**Correlation with human Likert score annotations:** We compute the Spearman and Kendall rank correlations between the human-annotated Likert scores and the metric values. These values were gathered for the 300 CodeReviewer test instances mentioned in section 4.1 The results are shown in Table 2. Notably, we exclude human annotations done on the CodeReviewer “ground truth” references (“Ground Truth” row in Table 9) because the reference-based metric value for these would be 1 by default unfairly lowering the correlation values for reference-based metrics that we compare against. We observe that while all the reference-based metrics have weak to no correlations with human judgment our metric *Rel* is the most correlated, achieving a moderate correlation with human annotations. We also show correlations between the human Likert scale annotations for all the dimensions and the reference-based metrics as well as the rest of our proposed metrics — *Con* and *Comp* in Table 8.

**Comparing system rankings:** Arguably the most important quality a metric should have is the ability to rank systems similar to human evaluators. To see if our metric *Rel* is capable of doing that we compare the system rankings produced by the human annotations (mean relevance Likert scores or “Human Annotations - Rel” column in Table 9) and by *Rel* (“Our Metric - Rel” column in Table 9), for the

9 evaluated systems. The system rankings (shown in Table 14) reveal that our metric gets the ranking mostly right, except for LLaMA-3-8B-Instruct which is ranked slightly lower by our metric compared to the human annotations. Additionally, we also compute the Spearman and Kendall rank correlations between system rankings according to each metric and the human annotations. Only our metric achieves a strong, statistically significant correlation with the human annotations.

We also compare the exact quality dimension scores for each system according to our metrics (*Con*, *Comp*, and *Rel*) and human annotations in Table 9. We normalize the human annotated Likert scores per dimension (*DimVal*) from 1 to 5 to 0 to 1 (*NormDimVal*) as:  $NormDimVal = \frac{DimVal-1}{4}$ . The table shows that our metric agrees with the human annotations over the best and worst systems. Our metric also has a similar spread of values to human annotations and exhibits greater variability between systems and sensitivity to review quality scores as shown in Figure 8 and Table 12.

### 5.3 CodeReviewer Dataset Reference Quality

We compare the quality of the CodeReviewer reference reviews with the reviews generated by the 9 system evaluated by the human annotators. The average scores for conciseness, comprehensiveness, and relevance attained by the CodeReviewer references are 3.05, 1.88, and 2.13, while the average scores obtained by all 9 systems are 2.57, 1.84, and 1.99. This suggests that the average reference review is barely better than the average evaluated system according to the human annotators for relevance, but they are more concise. Additionally the best system according to human annotators (as evident from Table 9), GPT-3.5 achieves average scores of 3.63, 2.65, and 2.9 for each dimension – much better than the reference reviews. This provides further motivation for the development of reference-free evaluation metrics like CRScore for code review.

### 5.4 Failure Cases

We analyze the cases where our metric greatly overestimates or underestimates the quality of a review with respect to human annotations. We find such cases using the procedure described in Appendix E.1.

For underestimation cases we observe our pseudo-reference generation pipeline generates fewer references on average (2.44) compared to

the whole data (4.76). This suggests that having fewer references makes it harder to evaluate the relevance of reviews. Additionally, we observe reviews like “Why do we need these imports” which are brief, contain stopwords, and have fewer relevant tokens, making it hard for STS to recognize their relevance to pseudo-reference.

For overestimation cases, we observe the presence of inline code snippets at a higher rate (45%) compared to all reviews (28%) and underestimation cases (12%). Some example reviews for both cases are shown in Table 15 and 16.

## 6 Discussion

In this work we identify issues with current code review evaluation benchmarks like CodeReviewer (Li et al., 2022) which fail to capture the one-to-many nature of code review and contain noisy references. To enable auditing current evaluation metrics and aid the development of a better metric we propose three review quality dimensions – conciseness, comprehensiveness, and relevance based on current literature on reference-free evaluation (Mehri and Eskenazi, 2020b; Piorkowski et al., 2020; Turzo and Bosu, 2024).

To ground these dimensions into topics that reviews should address (Rasheed et al., 2024) we propose an automated pseudo-reference generation pipeline that leverages Large Language Models (LLMs) and code smell detectors to generate claims and detect issues in code changes to be covered by reviews. We validate the quality of these pseudo-references via human evaluation. Based on these dimensions and pseudo-references we develop reliable guidelines for coding the quality of reviews and collect annotations for 9 review generation systems and the “ground truth” reviews for the CodeReviewer dataset spanning Python, Java and Javascript.

The collected annotations show that current reference based metrics indeed fail to capture human preferences which is further compounded by humans preferring some models over the references. We propose CRScore as a metric to capture the three dimensions using the pseudo-references through STS models like sentence transformers (Reimers, 2019) and show that our approach has the greatest alignment with human preferences in terms of review level scores (0.4577  $\tau$  and 0.5425  $r_s$ ) and system ranking (0.95  $r_s$  and 0.8889  $\tau$ ) and greatest sensitivity (Figure 8). Despite this we



note that the correlations are only moderate and CRScore has some issues with overestimating and underestimating review quality when reviews include inline code or our pseudo-reference pipeline generates very few claims, suggesting that using better STS matching methods and pseudo-reference generation methods can improve our metric.

## 7 Conclusion

Our work takes the first steps towards addressing the variety of challenges involved in evaluating the quality of code reviews. We propose useful dimensions for capturing review quality, propose a neuro-symbolic method for operationalizing them in a grounded way using automatically generated pseudo-references, collect a dataset of human judgment of review quality, and propose a reference-free automated metric called CRScore to measure review quality along the proposed dimensions using pseudo-references. We benchmark our metric against 7 reference-based metrics, for evaluating 9 diverse review generation systems over the human annotations and show that our metric achieves the best alignment and is the most sensitive, but there is still scope for improvement by developing better pseudo-reference generation and STS matching methods.

## 8 Future Work

As discussed before, although our metric takes a great first step towards better reference free evaluation of code reviews, it still suffers from systematic under and over estimation errors in certain cases. According to us these limitations ultimately stem from the lack of exhaustiveness of the pseudo-reference generation pipeline and the limitations of STS methods when it comes to matching data containing both code and text. We think research is required to add components that can detect issues of code security, efficiency, etc. to the pseudo-reference generation pipeline for greater coverage of possible issues and better embedding models should be developed to capture functional properties of code to improve STS methods for matching code and text.

Additional limitations of our study include coverage of only Python, Java, and Javascript code for annotations and code smell detection even though the CodeReviewer dataset contains more languages like C, C++, Go, and Ruby. Future work could explore ways to extend the code smell detection

and collect more annotations for these other languages.

## Limitations

- While annotating reviews for the review quality dimensions of conciseness, comprehensiveness and relevance, the human annotators are encouraged to use claims and issues that start off as the same list of pseudo-references used by CRScore. However we argue that this isn't a source of anchoring bias because the human annotators are allowed to add and remove claims from the pseudo-references while they annotate them from review quality. This means the final list of claims used for review quality annotations is different from the one used by CRScore. Additionally having a common list of claims and issues for the two annotators helped us improve the reliability of our coding framework as talked about in section 4.2. It also allowed the annotators to ground the comprehensiveness of reviews into a concrete set of things that a review should cover.
- Semantic textual similarity (STS) models are imperfect at matching relevant pseudo-references to the review sentences, especially when there are very few claims and the reviews contain inline code snippets, where the latter can inflate the STS scores.
- Our pseudo-reference generation pipeline is not comprehensive enough in some cases which can lead to underestimation of review quality as shown by our failure case analysis. Additionally, it could be extended by adding more modules like code smell detectors for aspects like code security, code efficiency, etc. similar to [Rasheed et al. \(2024\)](#). Also, code smell detector tools can be added for languages other than Python, Java, and Javascript, like Go, C/C++, and Ruby present in the CodeReviewer dataset.
- Our metric only achieves a moderate correlation with human annotations of review quality, and while it is much better than the reference-based metrics in terms of alignment with human judgment and sensitivity to review quality as judged by humans, it is only a first step towards developing better metrics for code review. Future work should try to address the

limitations of our claim generation pipeline and STS methods.

## Ethics Statement

We believe our work doesn't violate any ethical guidelines and is compliant with copyright rules and regulations as we use an existing publicly available dataset and augment it with annotations of review quality using reviews generated by 9 systems and the references in the dataset. While there is a slight risk of harmful or toxic text being a part of the pseudo-references generated by the LLM component in our pseudo-reference generation pipeline we don't believe it to be a major risk based on the annotations done for the pseudo-reference quality.

## Acknowledgements

This research was funded in part by NSF grant DSES 2222762.

## References

- [Apache lucene](#).
- Mouna Abidi, Manel Grichi, Foutse Khomh, and Yann-Gaël Guéhéneuc. 2019. [Code smells for multi-language systems](#). In *Proceedings of the 24th European Conference on Pattern Languages of Programs*, EuroPLop '19, New York, NY, USA. Association for Computing Machinery.
- AI@Meta. 2024. [Llama 3 model card](#).
- Deepika Badampudi, Michael Unterkalmsteiner, and Ricardo Britto. 2023. [Modern code reviews—survey of literature and practice](#). *ACM Transactions on Software Engineering and Methodology*, 32(4):1–61.
- Dzmitry Bahdanau. 2014. Neural machine translation by jointly learning to align and translate. *arXiv preprint arXiv:1409.0473*.
- Ramakrishna Bairi, Atharv Sonwane, Aditya Kanade, Arun Iyer, Suresh Parthasarathy, Sriram Rajamani, B Ashok, and Shashank Shet. 2024. Codeplan: Repository-level coding using llms and planning. *Proceedings of the ACM on Software Engineering*, 1(FSE):675–698.
- Gabriele Bavota and Barbara Russo. 2015. [Four eyes are better than two: On the impact of code reviews on software quality](#). In *2015 IEEE International Conference on Software Maintenance and Evolution (ICSME)*, pages 81–90.
- Moritz Beller, Alberto Bacchelli, Andy Zaidman, and Elmar Juergens. 2014. [Modern code reviews in open-source projects: Which problems do they fix?](#) In *Proceedings of the 11th Working Conference on Mining Software Repositories*, MSR 2014, page 202–211, New York, NY, USA. Association for Computing Machinery.
- Wikipedia Contributors. 2019a. [Cohesion \(computer science\)](#).
- Wikipedia Contributors. 2019b. [Cyclomatic complexity](#).
- Lee Dong-Kyu. 2024. A gpt-based code review system for programming language learning. *arXiv preprint arXiv:2407.04722*.
- Aryaz Eghbali and Michael Pradel. 2023. [Crystalbleu: Precisely and efficiently measuring the similarity of code](#). In *Proceedings of the 37th IEEE/ACM International Conference on Automated Software Engineering*, ASE '22, New York, NY, USA. Association for Computing Machinery.
- Mikhail Evtikhiev, Egor Bogomolov, Yaroslav Sokolov, and Timofey Bryksin. 2023. Out of the bleu:: How should we assess quality of the code generation models?
- Lishui Fan, Jiakun Liu, Zhongxin Liu, David Lo, Xin Xia, and Shanping Li. 2024. Exploring the capabilities of llms for code change related tasks. *arXiv preprint arXiv:2407.02824*.
- Zhangyin Feng, Daya Guo, Duyu Tang, Nan Duan, Xiaocheng Feng, Ming Gong, Linjun Shou, Bing Qin, Ting Liu, Daxin Jiang, and Ming Zhou. 2020. [Codebert: A pre-trained model for programming and natural languages](#). In *Findings of the Association for Computational Linguistics: EMNLP 2020, Online Event, 16-20 November 2020*, volume EMNLP 2020 of *Findings of ACL*, pages 1536–1547. Association for Computational Linguistics.
- Martin Fowler. 1997. Refactoring: Improving the design of existing code.
- Alexander Frömmgen, Jacob Austin, Peter Choy, Nimesh Ghelani, Lera Kharatyan, Gabriela Surita, Elena Khrapko, Pascal Lamblin, Pierre-Antoine Manzagol, Marcus Revaj, Maxim Tabachnyk, Daniel Tarlow, Kevin Vilella, Dan Zheng, Satish Chandra, and Petros Maniatis. 2024. Resolving code review comments with machine learning. In *2024 IEEE/ACM 46th International Conference on Software Engineering: Software Engineering in Practice (ICSE-SEIP)*.
- Kavita Ganesan. 2018. Rouge 2.0: Updated and improved measures for evaluation of summarization tasks. *arXiv preprint arXiv:1803.01937*.
- Daya Guo, Qihao Zhu, Dejian Yang, Zhenda Xie, Kai Dong, Wentao Zhang, Guanting Chen, Xiao Bi, Y Wu, YK Li, et al. 2024. Deepseek-coder: When the large language model meets programming-the rise of code intelligence. *CoRR*.

- Michael Hanna and Ondřej Bojar. 2021. [A fine-grained analysis of BERTScore](#). In *Proceedings of the Sixth Conference on Machine Translation*, pages 507–517, Online. Association for Computational Linguistics.
- Jack Hessel, Ari Holtzman, Maxwell Forbes, Ronan Le Bras, and Yejin Choi. 2021. [CLIPScore: A reference-free evaluation metric for image captioning](#). In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 7514–7528, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.
- Anwen Hu, Shizhe Chen, Liang Zhang, and Qin Jin. 2023. Infometic: An informative metric for reference-free image caption evaluation. *arXiv preprint arXiv:2305.06002*.
- Klaus Krippendorff. 2018. *Content analysis: An introduction to its methodology*. Sage publications.
- J R Landis and G G Koch. 1977. The measurement of observer agreement for categorical data. *Biometrics*, 33(1):159–174.
- Zhen Li, Xiaohan Xu, Tao Shen, Can Xu, Jia-Chen Gu, and Chongyang Tao. 2024. Leveraging large language models for nlq evaluation: A survey. *arXiv preprint arXiv:2401.07103*.
- Zhiyu Li, Shuai Lu, Daya Guo, Nan Duan, Shailesh Jannu, Grant Jenks, Deep Majumder, Jared Green, Alexey Svyatkovskiy, Shengyu Fu, et al. 2022. Automating code review activities by large-scale pre-training. In *Proceedings of the 30th ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering*, pages 1035–1047.
- Hong Yi Lin, Patanamon Thongtanunam, Christoph Treude, and Wachiraphan Charoenwet. 2024. Improving automated code reviews: Learning from experience. In *Proceedings of the 21st International Conference on Mining Software Repositories*, pages 278–283.
- Haiyang Liu, Yang Zhang, Vidya Saikrishna, Quanquan Tian, and Kun Zheng. 2024. Prompt learning for multi-label code smell detection: A promising approach. *arXiv preprint arXiv:2402.10398*.
- Xinghua Liu and Cheng Zhang. 2017. The detection of code smell on software development: a mapping study. In *2017 5th International Conference on Machinery, Materials and Computing Technology (ICMMCT 2017)*, pages 560–575. Atlantis Press.
- Junyi Lu, Lei Yu, Xiaojia Li, Li Yang, and Chun Zuo. 2023. Llama-reviewer: Advancing code review automation with large language models through parameter-efficient fine-tuning. In *2023 IEEE 34th International Symposium on Software Reliability Engineering (ISSRE)*, pages 647–658. IEEE Computer Society.
- Pranava Madhyastha, Josiah Wang, and Lucia Specia. 2019. [VIFIDEL: Evaluating the visual fidelity of image descriptions](#). In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 6539–6550, Florence, Italy. Association for Computational Linguistics.
- Shane McIntosh, Yasutaka Kamei, Bram Adams, and Ahmed E. Hassan. 2014. [The impact of code review coverage and code review participation on software quality: A case study of the qt, vtk, and itk projects](#). In *Proceedings of the 11th Working Conference on Mining Software Repositories, MSR 2014*, page 192–201, New York, NY, USA. Association for Computing Machinery.
- Shikib Mehri and Maxine Eskenazi. 2020a. Unsupervised evaluation of interactive dialog with dialogpt. In *Proceedings of the 21th Annual Meeting of the Special Interest Group on Discourse and Dialogue*, pages 225–235.
- Shikib Mehri and Maxine Eskenazi. 2020b. [USR: An unsupervised and reference free evaluation metric for dialog generation](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 681–707, Online. Association for Computational Linguistics.
- mschwager. 2016. [Github - mschwager/cohesion: A tool for measuring python class cohesion](#).
- Niklas Muennighoff, Nouamane Tazi, Loïc Magne, and Nils Reimers. 2022. [Mteb: Massive text embedding benchmark](#). *arXiv preprint arXiv:2210.07316*.
- Thanis Paiva, Amanda Damasceno, Eduardo Figueiredo, and Cláudio Sant’Anna. 2017. [On the evaluation of code smells and detection tools](#). *Journal of Software Engineering Research and Development*, 5(1):7.
- Arjun Panickssery, Samuel R Bowman, and Shi Feng. 2024. Llm evaluators recognize and favor their own generations. *arXiv preprint arXiv:2404.13076*.
- Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. [Bleu: a method for automatic evaluation of machine translation](#). In *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, pages 311–318, Philadelphia, Pennsylvania, USA. Association for Computational Linguistics.
- Nikhil Pinnaparaju, Reshith Adithyan, Duy Phung, Jonathan Tow, James Baicoianu, and Nathan Cooper. [Stable code 3b](#).
- David Piorkowski, Daniel González, John Richards, and Stephanie Houde. 2020. Towards evaluating and eliciting high-quality documentation for intelligent systems. *arXiv preprint arXiv:2011.08774*.
- Maja Popović. 2015. [chrF: character n-gram F-score for automatic MT evaluation](#). In *Proceedings of the Tenth Workshop on Statistical Machine Translation*, pages 392–395, Lisbon, Portugal. Association for Computational Linguistics.

- Maja Popović. 2017. [chrF++: words helping character n-grams](#). In *Proceedings of the Second Conference on Machine Translation*, pages 612–618, Copenhagen, Denmark. Association for Computational Linguistics.
- Chanathip Pornprasit and Chakkrit Tantithamthavorn. 2024. [Fine-tuning and prompt engineering for large language models-based code review automation](#). In *Information and Software Technology*, page 107523.
- Md Tajmilur Rahman, Rahul Singh, and Mir Yousuf Sultan. 2024. Automating patch set generation from code reviews using large language models. In *Proceedings of the IEEE/ACM 3rd International Conference on AI Engineering-Software Engineering for AI*, pages 273–274.
- Zeeshan Rasheed, Malik Abdul Sami, Muhammad Waseem, Kai-Kristian Kemell, Xiaofeng Wang, Anh Nguyen, Kari Systä, and Pekka Abrahamsson. 2024. Ai-powered code review with llms: Early results. *arXiv preprint arXiv:2404.18496*.
- N Reimers. 2019. Sentence-bert: Sentence embeddings using siamese bert-networks. *arXiv preprint arXiv:1908.10084*.
- Ehud Reiter. 2018. [A structured review of the validity of BLEU](#). *Computational Linguistics*, 44(3):393–401.
- Shuo Ren, Daya Guo, Shuai Lu, Long Zhou, Shujie Liu, Duyu Tang, M. Zhou, Ambrosio Blanco, and Shuai Ma. 2020. [Codebleu: a method for automatic evaluation of code synthesis](#). *ArXiv*, abs/2009.10297.
- Baptiste Roziere, Jonas Gehring, Fabian Gloeckle, Sten Sootla, Itai Gat, Xiaoqing Ellen Tan, Yossi Adi, Jingyu Liu, Tal Remez, Jérémy Rapin, et al. 2023. Code llama: Open foundation models for code. *arXiv preprint arXiv:2308.12950*.
- Rana Sandouka and Hamoud Aljamaan. 2023. Python code smells detection using conventional machine learning models. *PeerJ Comput Sci*, 9:e1370.
- Tushar Sharma, Vasiliki Efstathiou, Panos Louridas, and Diomidis Spinellis. 2021. [Code smell detection by deep direct-learning and transfer-learning](#). *Journal of Systems and Software*, 176:110936.
- Ngoc Tran, Hieu Tran, Son Nguyen, Hoan Nguyen, and Tien Nguyen. 2019. Does bleu score work for code migration? In *2019 IEEE/ACM 27th International Conference on Program Comprehension (ICPC)*, pages 165–176. IEEE.
- Nikolaos Tsantalis, Theodoros Chaikalis, and Alexander Chatzigeorgiou. 2008. [Jdeodorant: Identification and removal of type-checking bad smells](#). In *2008 12th European Conference on Software Maintenance and Reengineering*, pages 329–331.
- Rosalia Tufano, Simone Masiero, Antonio Mastropaolo, Luca Pascarella, Denys Poshyvanyk, and Gabriele Bavota. 2022. Using pre-trained models to boost code review automation. In *Proceedings of the 44th international conference on software engineering*, pages 2291–2302.
- Rosalia Tufano, Luca Pascarella, Michele Tufano, Denys Poshyvanyk, and Gabriele Bavota. 2021. Towards automating code review activities. In *2021 IEEE/ACM 43rd International Conference on Software Engineering (ICSE)*, pages 163–174. IEEE.
- Asif Kamal Turzo and Amiangshu Bosu. 2024. What makes a code review useful to opendev developers? an empirical investigation. *Empirical Software Engineering*, 29(1):6.
- Manushree Vijayvergiya, Małgorzata Salawa, Ivan Budiselić, Dan Zheng, Pascal Lamblin, Marko Ivanković, Juanjo Carin, Mateusz Lewko, Jovan Andonov, Goran Petrović, et al. 2024. Ai-assisted assessment of coding practices in modern code review. *arXiv preprint arXiv:2405.13565*.
- Yuxiang Wei, Zhe Wang, Jiawei Liu, Yifeng Ding, and Lingming Zhang. 2023. Magicoder: Source code is all you need. *arXiv preprint arXiv:2312.02120*.
- Zezhou Yang, Cuiyun Gao, Zhaoqiang Guo, Zhenhao Li, Kui Liu, Xin Xia, and Yuming Zhou. 2024. A survey on modern code review: Progresses, challenges and opportunities. *arXiv preprint arXiv:2405.18216*.
- Yongda Yu, Guoping Rong, Haifeng Shen, He Zhang, Dong Shao, Min Wang, Zhao Wei, Yong Xu, and Juhong Wang. 2024. [Fine-tuning large language models to improve accuracy and comprehensibility of automated code review](#). *ACM Trans. Softw. Eng. Methodol.* Just Accepted.
- Fengji Zhang, Zexian Zhang, Jacky Wai Keung, Xiangru Tang, Zhen Yang, Xiao Yu, and Wenhua Hu. 2024. [Data preparation for deep learning based code smell detection: A systematic literature review](#).
- Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q. Weinberger, and Yoav Artzi. 2020. [Bertscore: Evaluating text generation with BERT](#). In *8th International Conference on Learning Representations, ICLR 2020, Addis Ababa, Ethiopia, April 26-30, 2020*. OpenReview.net.
- Shuyan Zhou, Uri Alon, Sumit Agarwal, and Graham Neubig. 2023. [Codebertscore: Evaluating code generation with pretrained models of code](#). *Deep Learning for Code (DL4C) workshop, 11th International Conference on Learning Representations, ICLR 2023*, abs/2302.05527.

## A Introduction Details

This section contains details to corroborate some of the points made in the introduction section.



### A.1 Benefits of Neuro Symbolic Pseudo-Reference Generation

While LLMs have recently shown promise for evaluating natural language generation (NLG) (Li et al., 2024) they suffer from biases like favoring their own generations (“self-selection bias”) (Panicssery et al., 2024) or in other words if we were to have Magicoder or GPT-3.5 as the evaluator LLM it would assign higher scores to text generated by Magicoder and GPT-3.5 respectively. Code analysis tools (CATs) on the other hand are limited in scope compared to LLMs in detecting issues like best practice violations (Vijayvergiya et al., 2024) but don’t have any self-selection bias. However, combining these methods can reduce the self-selection bias of LLMs, while supplementing the narrow coverage of code analysis tools. Indeed the results show that despite using Magicoder as the evaluation LLM, our metric CRScore doesn’t preferentially rank Magicoder above any models other than LLaMA-3 when compared to the human ranking.

Review	Missing Context
Don’t redefine, just import the existing one in <code>cmdline.py</code> . :)	Folder structure, Codebase organization
I think we can remove this function, right? (duplicate with <code>ses_starter.py</code> )	Folder structure, Codebase organization
MPRester(os.environ(["MP_API_KEY"])) can be left simply as MPRester() and it will pick up the API key from the environment. What ...	Class definition, Environment variables

Table 4: Ground truth reviews in the automatically mined CodeReviewer data that assume contextual information about the code base not available in the dataset

## B More Related Work

### B.1 Code Specific Reference Based Metrics

Due to the popularity and convenience of automated reference-based metrics like BLEU (Papineni et al., 2002), ROUGE (Ganesan, 2018), and BERTScore (Zhang et al., 2020) the research community has developed several code-specific versions like CodeBLEU (Ren et al., 2020), RUBY

(Tran et al., 2019), CrystalBLEU (Eghbali and Pradel, 2023) and CodeBERTScore (Zhou et al., 2023). CodeBLEU extends BLEU by incorporating code structure through dataflow and syntax match between generated code and references, while CrystalBLEU filters out trivially shared n-grams. RUBY incorporates the distance between the syntax tree and program dependency graph of references and generated code. CodeBERTScore extends the embedding-based BERTScore by replacing BERT with a pre-trained CodeBERT (Feng et al., 2020) model. However prior studies have shown metrics like BLEU to have low validity (overlap with human judgment) and reliability for text generation (Reiter, 2018), code generation (Evtikhiev et al., 2023), and code migration (Tran et al., 2019). However, metrics like ROUGE, BERTScore, and CodeBERTScore all have a notion of precision, recall, and f-score which is captured by conciseness, comprehensiveness, and relevance respectively.

### B.2 Code Review Automation

Due to the high time and resource demands of code review automated approaches have gained popularity (Yang et al., 2024). Pornprasit and Tantihamthavorn (2024); Lu et al. (2023); Dong-Kyu (2024); Fan et al. (2024); Yu et al. (2024) propose fine-tuning and prompt engineering approaches to leverage LLMs for code review and code-change related tasks. Frömmgen et al. (2024); Rahman et al. (2024) propose methods for code refactoring based on review comments. Vijayvergiya et al. (2024) propose the detection of “best practice violations”. Lin et al. (2024) propose oversampling reviews from experienced reviewers as a proxy of review quality improving informativeness and correctness of generated reviews. Rasheed et al. (2024) propose an LLM agent for code review and code smell detection.

### B.3 Code Smell Detection

The problem of detecting “code smells” or symptoms of design flaws and bad practices (also called *anti-patterns*) has been traditionally tackled by analysis-based approaches (Tsantalis et al., 2008; Paiva et al., 2017; Liu and Zhang, 2017). Recent work has explored learning-based methods for potentially more nuanced detection of code smells. Sandouka and Aljamaan (2023) create a dataset of 1k Python code smells like “Long Method” and “Large Class” to train traditional ML models like random forests. Sharma et al. (2021) leverage

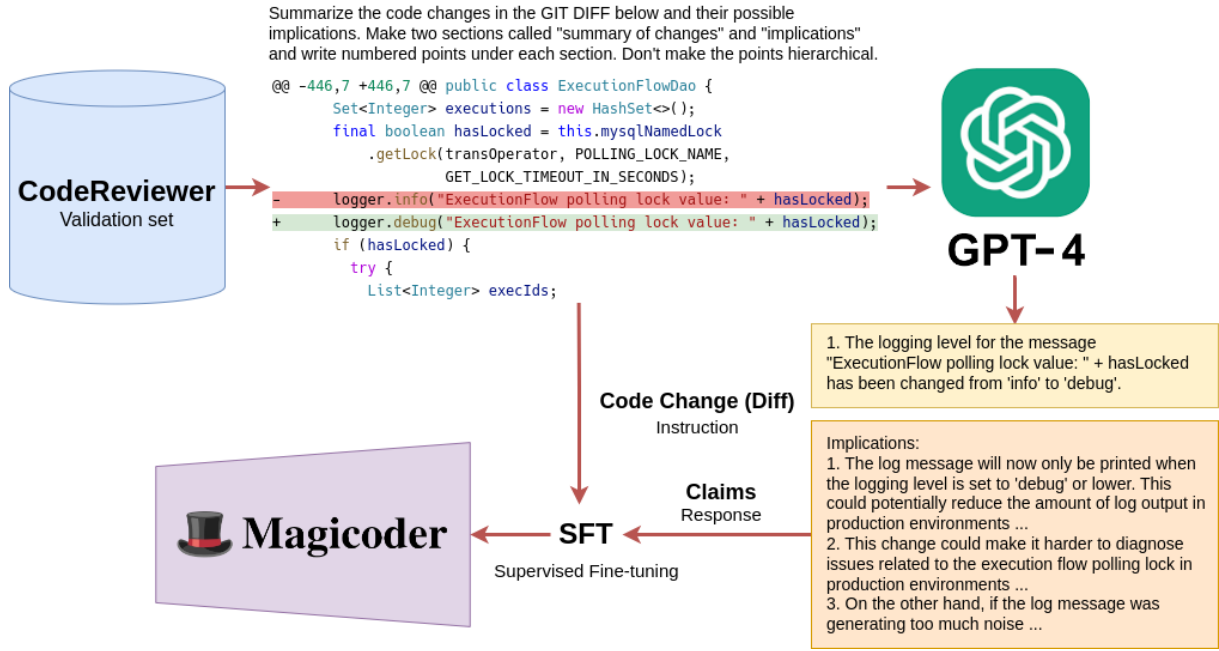


Figure 3: Supervised fine-tuning pipeline for training Magicoder-6.7B for claim generation. We generate synthetic data by using GPT-4 to generate claims for the code changes in CodeReviewer validation set.

deep learning models like autoencoders and transfer learning for adapting to unseen programming languages. Liu et al. (2024) propose a prompting-based approach for “Long Method” and “Long Parameter List” code smells in Java. Rasheed et al. (2024) create an LLM agent to detect code smells in repositories. However, these approaches are limited in the types of smells they target, training data requirements (Zhang et al., 2024), or lack comprehensive evaluation. Also, code smells differ across programming languages (Abidi et al., 2019), and transfer learning approaches can only be leveraged for similar languages (Sharma et al., 2021). Due to these limitations of learning-based methods we use traditional language-specific code analysis tools.

## C Method Details

This appendix contains additional details on the implementation of our CRScore metric.

### C.1 Distribution of Sentence Similarity Scores

We plot the histogram of values of the sentence similarity scores in Figure 5 showing a roughly normal distribution. We also plot the quantile-quantile (Q-Q) plot in Figure 6 that compares the quantile of a normal distribution with the empirically observed distribution of sentence similarity scores. Ideally, the Q-Q plot should be a straight line (shown in red) but we observe deviation to-

wards really high ranges among the actual values (shown in blue). Due to computational constraints, these plots are constructed out of a randomly sampled subset of 100k similarity scores from the original 100M+ sentence similarity scores computed from the CodeReviewer test set review pairs.

### C.2 Code Smell Detection Details

In this section, we cover some of the details of the code smell detectors used in this study.

#### C.2.1 Class Cohesion

Class cohesion captures the degree to which the elements of a class belong together (Contributors, 2019a). In other words, cohesion measures the strength of the relationship between pieces of functionality (attributes and methods) within a given class. For example, in highly cohesive classes functionality is strongly related and methods and attributes are more co-dependant and hang together as a logical whole (mschwager, 2016).

#### C.2.2 Cyclomatic Complexity

Cyclomatic complexity is a software metric used to indicate the complexity of a program (Contributors, 2019b). It is a quantitative measure of the number of linearly independent paths through a program’s source code.

Popular tools like Radon<sup>7</sup>, a cyclomatic com-

<sup>7</sup><https://pypi.org/project/radon/>

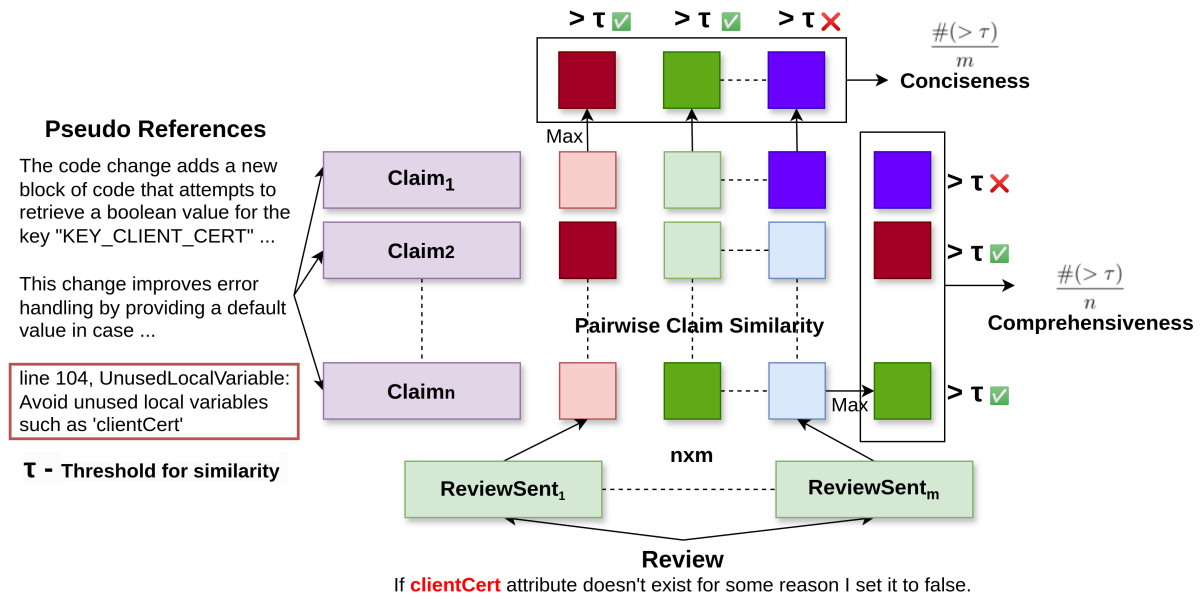


Figure 4: This figure shows how semantic textual similarity (STS) is used to measure the coverage of pseudo-references by the review sentences. We compute pairwise semantic similarities between all pseudo references and review sentences and employ a threshold to compute comprehensiveness as the fraction of pseudo references for which at least one review sentence has higher similarity than the threshold. Meanwhile, conciseness is the fraction of review sentences which high have higher similarity than the threshold with any pseudo reference.

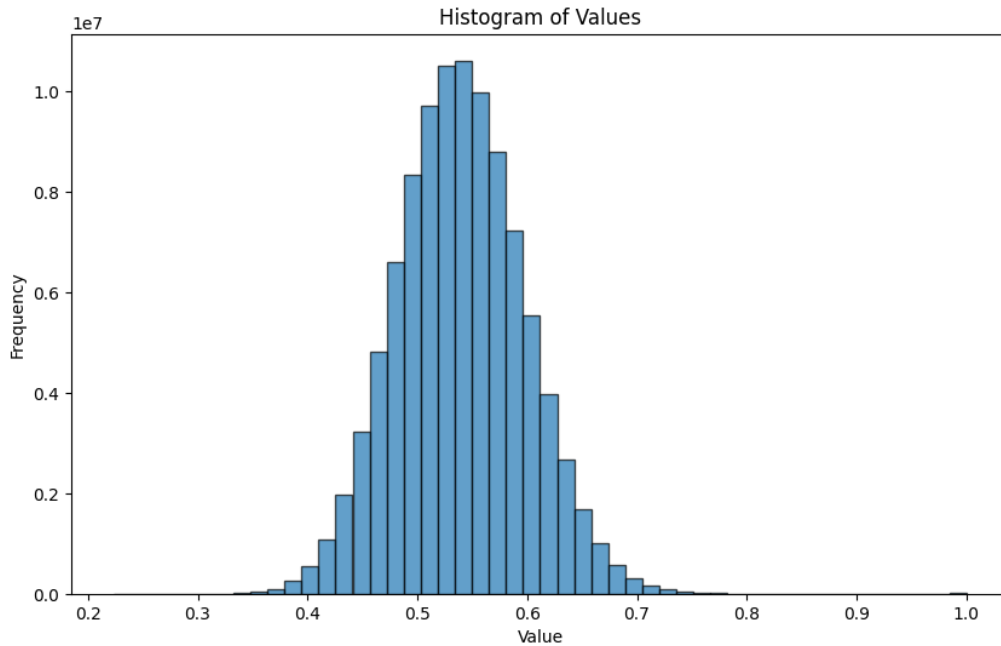


Figure 5: Histogram of sentence similarity of randomly sampled 100K sentence pairs from the CodeReviewer test set showing the scores are roughly normally distributed, justifying the usage of the 5-sigma rule for coming up with the threshold of 0.85 for high similarity used in metric computation.

plexity computation tool for Python often resort to a rank-based system that categorizes code blocks based on their cyclomatic complexity and associates them with letter grades as shown in Table 5.

Here “A” is the best grade and “F” is the worst grade. Code blocks are flagged for code smells if they have a complexity higher than or equal to “C” in Table 5.

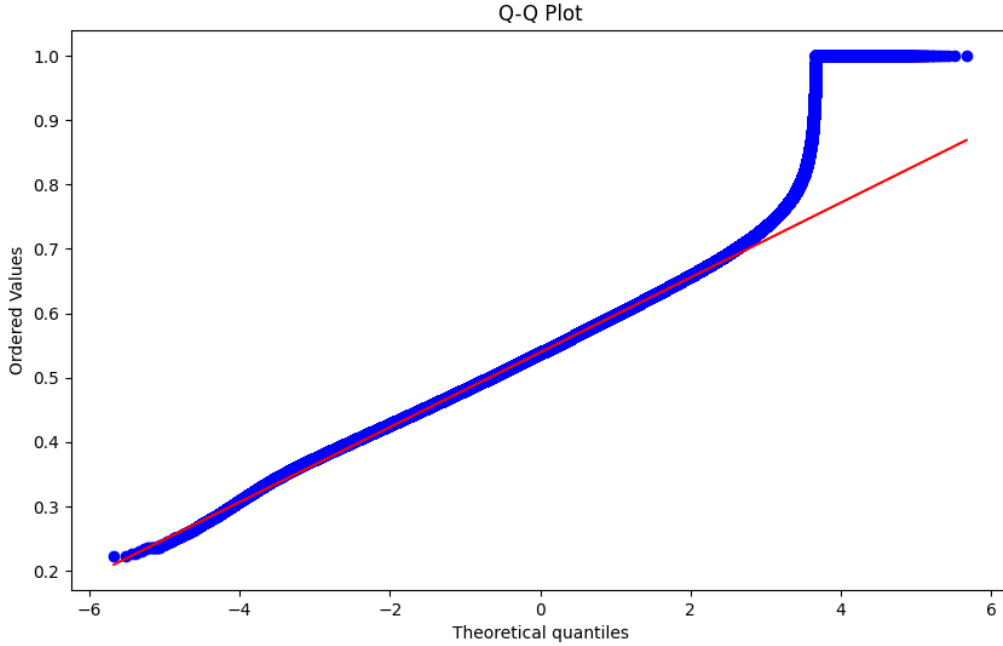


Figure 6: Q-Q plot comparing quantiles of empirically observed sentence similarity scores computed over 100K sentence pairs from the CodeReviewer test set showing the theoretical quantiles match a normal distribution except for really high values. The discrepancy seen here is likely due to the random sample being a smaller subset of the whole 100M+ sentence pairs for which we compute similarities.

CC Score	Rank	Risk
1-5	A	low - simple block
6-10	B	low - well structured and stable block
11-20	C	moderate - slightly complex block
21-30	D	more than moderate - more complex block
31-40	E	high - complex block, alarming
41+	F	very high - error-prone, unstable block

Table 5: Code complexity ranks, associated score thresholds, and descriptions of the potential risks.

## D Experimental Details

Further experimental details like guidelines for annotation of pseudo-reference accuracy, review quality, etc.

### D.1 Codebook for Rating Pseudo-Reference Quality

The raters are shown model-generated pseudo references (claims and implications) about the “diff”, which either describe the changes or speculate about potential implications. The claims are statements about the diff related to what changes took place, while the implications cover the effects of the changes or their interpretations like whether they implement a new functionality or even the

potential intent of the developers. Additionally, the raters are also given source files as context, including the versions of the file before and after the code change captured by the diff. The raters are expected to refer to them if they need more information than just the diff to judge the accuracy of the pseudo-references.

Given these inputs, the raters are supposed to code each pseudo-reference as shown in Table 7. In further analysis, we excluded the ambiguous claims (code 2) because they were rarely encountered.

### D.2 Rating Review Quality using Pseudo-References

The raters are given pseudo-references corresponding to the diff but unlike the pseudo-reference quality annotations, they also include issues/smells detected by the static analysis tools. The issues span formatting issues, bad programming practices, or more abstract patterns known as “code smells”. Code smells are heuristics or code characteristics associated with deeper problems concerning system design. It is important to note that they are not bugs, but rather subjective principles that vary across programming languages, developers, teams, etc. Given the diff and the pseudo-references, the



raters evaluate the quality of the review along the three dimensions — comprehensiveness, conciseness, and relevance. The raters are again given access to the source files for context and are asked to rate the quality of 10 reviews per code change that are generated by a diverse set of review generation systems and one of them is also the ground truth review from the CodeReviewer dataset.

The rater’s task is to first go through the list of pseudo-references and rate their necessity with respect to the code change. This stage is meant to remove any pseudo-references that are unnecessary or redundant with respect to the whole set. The second stage of the annotation is to link/associate various pseudo-references to each of the 10 reviews based on which of them are addressed in the review. Finally, the raters assign a score on a Likert scale of 1 to 5 to each review for each dimension. Some rules of thumb for assigning each score for each dimension are given in Table 6.

Raters are also given a helpful mental framework to help with the process of linking pseudo-references to the reviews. To explain what it means for a review to address a pseudo-reference, we give an analogy to aspect-based product reviews. The list/set of claims is similar to a list/set of product aspect descriptions, while the code review comment is similar to a product review. For example, for the product description and review below:

**Description of a newer version of a Phone:**

- The new phone improves battery life by 50% (battery life)
- The new screen has a higher resolution (screen)

**Review:** The new screen quality is great but the battery runs out quickly.

The review here is talking about both the screen and battery life aspects of the product so you can say it is addressing both aspects (or claims).

Now we can consider an example from the domain of code review:

**Pseudo-references for code change:**

- The code change is in the logging of errors in the response handler. (error logging)
- The formatting of the error message has been changed. (formatting change of error message)
- The previous formatting used '%s %' at the end of the error message, which was removed in the updated code. (% at the end of the error message)
- The change in the formatting of the error message will affect the way errors are logged and displayed. (effect of formatting on error logging)

**Code review:** The %s in the error message is redundant, and the indent=4 in json.dumps is unnecessary.

The review talks about the third claim by mentioning the %s style string at the end of the error logging.

### D.3 Simple Baseline Implementation Details

We create two simple baselines for code review generation a BM-25, kNN style retriever-based approach, and a seq2seq LSTM style model. The implementation details of both approaches are described below:

**BM-25 retriever:** The BM-25 retriever retrieves a relevant review by matching the closest code change from the CodeReviewer train set to the code change to be reviewed from the test set. For an efficient implementation, we create an inverted index from all the code changes in the train set using the Lucene searcher (`luc`) class from Pyserini<sup>8</sup>.

**LSTM reviewer:** We train a single hidden layer encoder-decoder seq2seq LSTM model with Bahdanau attention (Bahdanau, 2014) from scratch on the CodeReviewer training data. We train it with an Adam optimizer, and negative log-likelihood loss for about 10 epochs, saving the model with the least loss on the CodeReviewer validation set.

## E Additional Results and Analysis

### E.1 Finding Failure Cases of CRScore

We divide the scores spanned by our metric into 5 equally sized bins (Q1, Q2, Q3, and Q4 being the quantiles) and then find reviews where the metric underestimates the Rel value (value less than Q1 but a human rating of 5) and overestimates the Rel value (value greater than Q4 but a human rating of 1). We find 16 (0.61%) and 98 (3.74%) cases respectively for underestimation and overestimation, moreover, without these cases, our metric ( $\tau_{best}$ ) attains correlations of  $\tau = 0.5462$  and  $r_s = 0.6431$  and  $\tau = -0.5403$  and  $r_s = 0.6131$  for these cases, indicating their influence despite being less than 4% of the data.

<sup>8</sup><https://pypi.org/project/pyserini/>

Dimension	Score	Rule of thumb
Conciseness	1	none of the review is related to the claims+issues.
	2	some of the review is related to the claims+issues
	3	roughly half of the review is related to the claims+issues
	4	most of the review is related to the claims+issues
	5	basically the whole review is related to the claims+issues
Comprehensiveness	1	whole review doesn't cover any of the claims+issues
	2	review covers at least 1 claim or issue
	3	review covers roughly cover half the claims+issues
	4	review covers more than half/most of the claims+issues
	5	review covers practically all the claims+issues
Relevance		<p>relevance score must be between conciseness and comprehensiveness scores</p> <p>if one of the two dimensions is "limiting", i.e. has a very low score or poor quality for a given review then the relevance score should be biased/limited by it.</p> <p>E.g. a highly concise review but with very low comprehensiveness should have a relevance score close to the comprehensiveness</p>

Table 6: Some rules of thumb for scoring the quality of reviews for each dimension. These rules are meant as guidelines to calibrate multiple annotators and reduce the impact of learning effects in the early stages of annotation.

Code	Description	Examples	Explanation
1	Correct: You can find concrete evidence to validate or confirm a claim (either in the diff, the context/source files, or by looking up domain-specific knowledge on the web)	Any existing calls to the 'push' function will need to be updated to include the new 'hash' parameter. This could potentially break compatibility with older code.	This is true since if the parameters are passed by value and there is a parameter after the hash then its value will be accidentally passed as the hash
0	Incorrect: You can find concrete evidence to contradict a claim (either in the diff, the context/source files, or looking up domain-specific knowledge on the web)	The code changes involve the modification of the parameters passed to the ScalarSpaceEncoder function.	False because ScalarSpaceEncoder is a class and not a function.
-1	Unverifiable: You can't find evidence to confirm or contradict a claim (even after looking at the diff, source files, or looking up domain-specific knowledge on the web)	The changes could potentially affect the performance of the code as the order of the arguments does not matter any longer.	For this claim, the arguments were being passed by keyword so the order didn't matter for functional correctness but we can't comment on the performance
2	Ambiguous: The claim is underspecified and has multiple interpretations making it hard to determine what is to be tested/validated	The function will now only work with valid elements, preventing potential issues down the line.	It is underspecified (unclear) what "valid elements" means.

Table 7: Guidelines for coding accuracy of pseudo-references with example pseudo-references falling within each category and explanations for why they fall in that category.

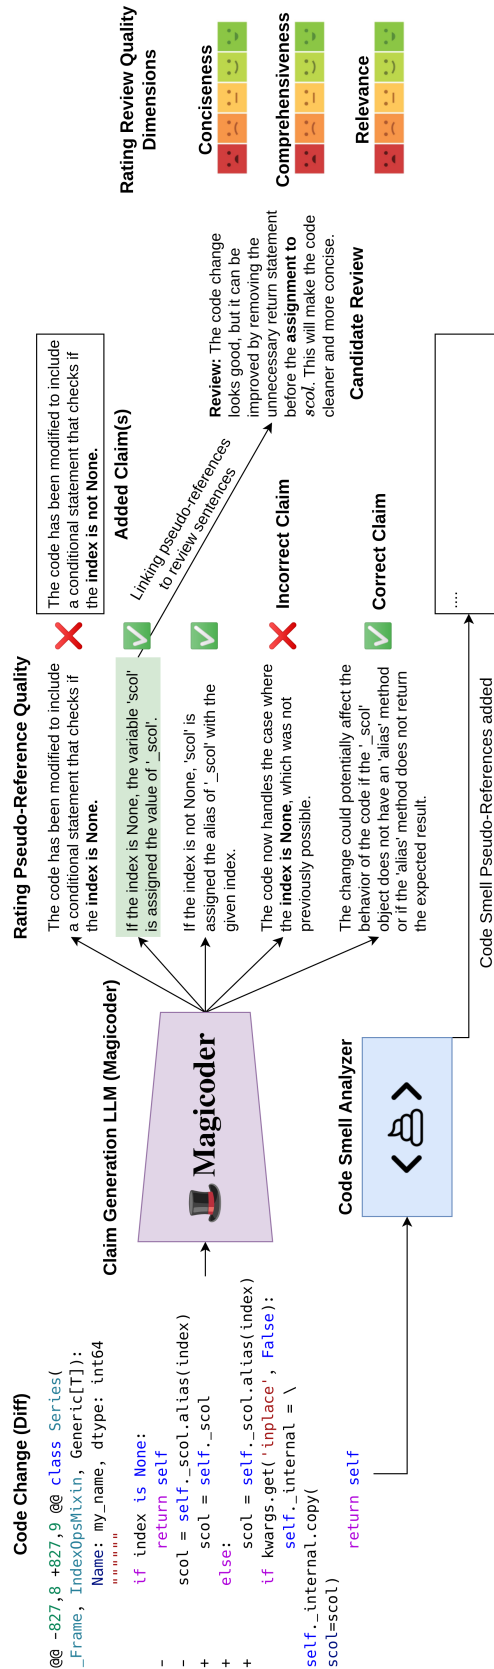


Figure 7: The stages of the annotation pipeline. For the first stage — Rating Pseudo-Reference Quality the annotators mark correct, incorrect, and unverifiable claims, while also adding any missing claims. In the second stage — Rating Review Quality Dimensions the annotators are given the updated set of pseudo-references from the first stage along with any pseudo-references generated by the code smell analysis tools to rate candidate reviews on the quality dimensions — conciseness, comprehensiveness, and relevance on a Likert scale using the pseudo-references.

Metric	Human Annotations					
	Con ( $\tau$ )	Comp ( $\tau$ )	Rel ( $\tau$ )	Con ( $r_s$ )	Comp ( $r_s$ )	Rel ( $r_s$ )
BLEU	0.0306	-0.0227	0.001	0.0358	-0.0293	-0.0001
BLEU (without stopwords)	0.0632	0.0221	0.0425	0.0776	0.0293	0.0542
BERTScore	0.1035	0.0622	0.081	0.1378	0.0813	1.083
Normalized Edit Distance	-0.0146	0.0443	0.0193	-0.0218	0.0584	0.0249
ROUGE-L	0.0921	0.0577	0.0757	0.1173	0.0758	0.0989
F-measure						
chrF	0.1236	0.1431	0.1484	0.1628	0.1874	0.1966
chrF++	0.1294	0.1496	0.1555	0.1707	0.1959	0.2057
Con ( $\tau_{GT}$ ) ours	0.4168			0.475		
Comp ( $\tau_{GT}$ ) ours		<b>0.4982</b>			<b>0.5832</b>	
Rel ( $\tau_{GT}$ ) ours			0.4437			<b>0.5405</b>
Con ( $\tau_{best}$ ) ours	<b>0.4491</b>			<b>0.5049</b>		
Comp ( $\tau_{best}$ ) ours		0.4974			0.5754	
Rel ( $\tau_{best}$ ) ours			<b>0.4567</b>			0.5431

Table 8: Kendall and Spearman rank correlations between human annotation for all dimensions: conciseness (Con), comprehensiveness (Comp), and relevance (Rel) and all metrics including our proposed Con, Comp and Rel metrics for both threshold values  $\tau_{GT}$  and  $\tau_{best}$ .

Model	Human Annotations			Our Metric		
	Con	Comp	Rel	Con	Comp	Rel
BM-25 retriever	<u>0.0301</u>	<u>0.0112</u>	<u>0.0163</u>	<u>0</u>	<u>0</u>	<u>0</u>
LSTM	0.1048	0.0361	0.0515	0.0372	0.0123	0.0179
CodeReviewer	0.5146	0.1692	0.2311	0.4639	0.2413	0.2974
Stable-Code-3B	0.3222	0.1383	0.1718	0.3353	0.2091	0.2319
DeepSeekCoder-6.7B-Instruct	0.6108	0.3153	0.3797	0.5037	0.3989	0.4043
Magocoder-S-DS-6.7B	0.4915	0.3127	0.3351	0.4381	0.4746	0.4032
LLaMA-3-8B-Instruct	0.6091	0.3368	0.4046	0.3503	0.3967	0.3404
CodeLLaMA-13B	0.1985	0.1546	0.1564	0.2493	0.2528	0.2309
GPT-3.5	<b>0.6564</b>	<b>0.4132</b>	<b>0.4759</b>	<b>0.5622</b>	<b>0.6301</b>	<b>0.5507</b>
Ground Truth	0.5129	0.219	0.2819	0.251	0.1598	0.1741

Table 9: Comparison of our proposed metrics with human annotations for the review quality dimensions. The results are reported on the subset of 300 annotated CodeReviewer test instances. We highlight the best-performing model (bold) and the worst-performing model (underlined) according to the human annotations and our metric.



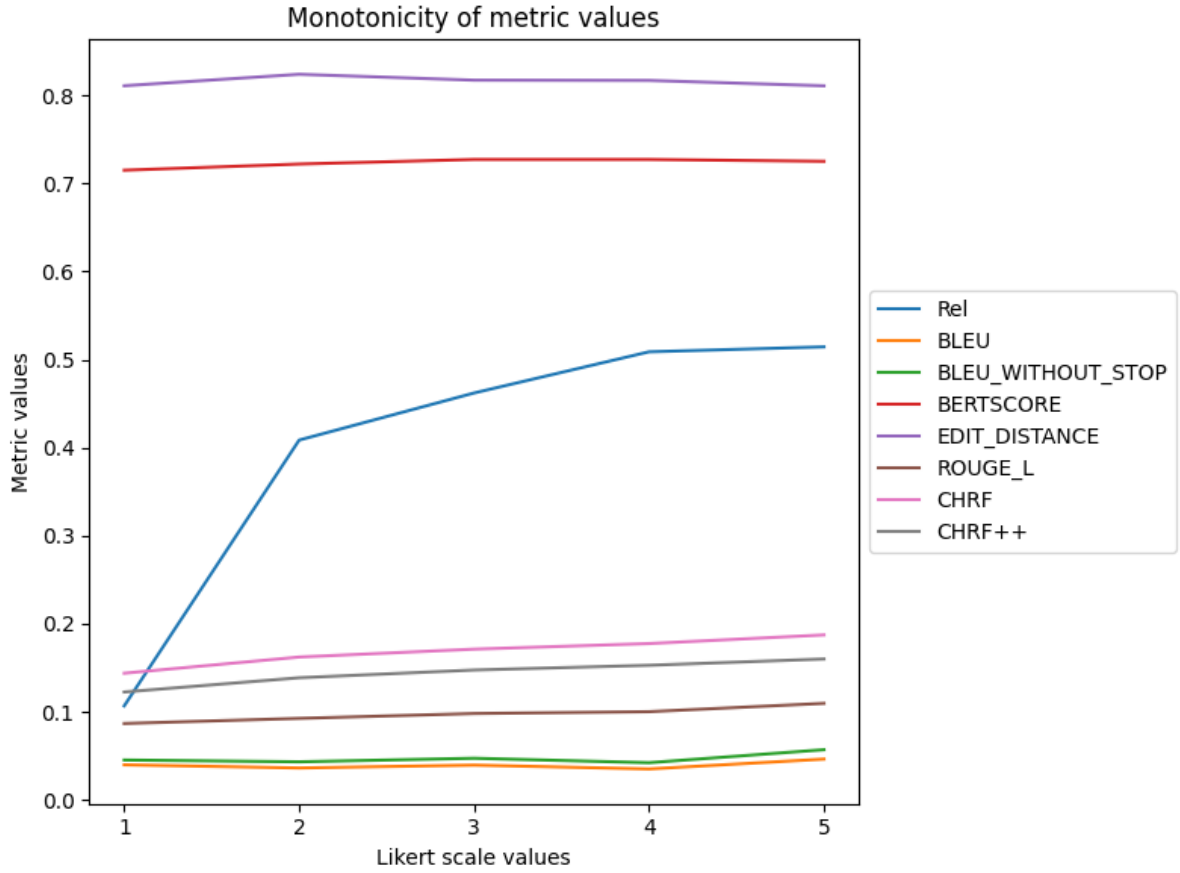


Figure 8: Monotonicity of metric values for reviews of various quality based on the Likert scale human annotations. The Rel metric exhibits the most variation across reviews of different quality while the other metrics have flat plots indicating that they fail to distinguish between reviews of varying quality meaningfully

Metric	Python		Java		Javascript	
	$\tau$	$r_s$	$\tau$	$r_s$	$\tau$	$r_s$
BLEU	-0.0017	-0.0006	-0.0101	-0.015	0.007	0.009
BLEU (without stopwords)	0.0532	0.0665	0.035	0.0442	0.0313	0.0405
BERTScore	0.0696	0.0942	0.0808	0.1081	0.0846	0.1128
Normalized Edit Distance	0.0367	0.046	0.0161	0.0217	0.0109	0.0129
ROUGE-L	0.0935	0.1201	0.0553	0.0724	0.075	0.0981
F-measure	0.1309	0.1746	0.1573	0.2077	0.1497	0.1982
chrF	0.1387	0.1848	0.1676	0.2208	0.1531	0.2021
chrF++	0.1387	0.1848	0.1676	0.2208	0.1531	0.2021
<b>Rel (<math>\tau_{GT}</math>) (Ours)</b>	0.4742	0.5788	0.3889	0.4738	0.473	0.5738
<b>Rel (<math>\tau_{best}</math>) (Ours)</b>	0.4746	0.5666	0.4093	0.4849	0.4904	0.5816

Table 10: Comparing Kendall-Tau ( $\tau$ ) and Spearman Rank ( $r_s$ ) correlation of reference-based evaluation metrics across each language annotated (Python, Java and Javascript) and our reference-free relevance score (Rel (F)) with human annotations for relevance. Correlations that are not statistically significant (p-value < 0.05) are grayed out.

Error Type	Description	Incorrect Reference	Corrected Reference	Frequency (%)
Knowledge Error	Pseudo-reference exhibits incorrect domain knowledge	The 'optparse' module is being imported with a comment indicating that it is being disabled due to its deprecation.	The 'optparse' module is being imported with a comment disabling a pylint deprecated-module warning	6.52
Reasoning Error	Pseudo-reference exhibits wrong logic applied by the pseudo-reference generator	Now, the 'can_edit_record' variable is only true if the function 'check_user_can_edit_record' returns true and the 'format' variable does not start with 't' (in lowercase).	Now, the 'can_edit_record' variable is only true if the function 'check_user_can_edit_record' returns true and the 'format' variable does not start with 't' (in any lower or uppercase).	10.87
Localization Error	The pseudo-reference generator misunderstands where a code change has taken place	The assertion in the test method "test_idxmapping_add_dimension" has been modified.	The assertion in the test method "test_idxmapping_redim" has been modified.	4.35
Over-generalization Error	The pseudo-reference generator makes an incorrect assumption/generalization from the code change	The addition of these import statements suggests that the code in this file will now be using the ResLayer and SimplifiedBasicBlock classes from the mmdet.models.utils package.		19.57
Comprehension Error	The pseudo-reference seems to "misread" the code change (like thinking removed lines are added, etc.)	The import statement for 'filter', 'range', and 'zip' has been moved from 'scapy.modules.six.moves' to 'scapy.modules.six'.	'filter' is now also imported from 'scapy.modules.six.moves'	58.7

Table 11: The various types of errors identified, their descriptions and examples (pseudo-references before and after correction of the error are shown) as well as relative frequencies as percentages are shown here. For this analysis, we annotated 46 erroneous pseudo-references

Model	BLEU	BLEU (without stop)	BERTScore	Norm. Edit Distance	ROUGE L f-score	chrF	chrF++	Con ( $\tau_{best}$ )	Comp ( $\tau_{best}$ )	Rel ( $\tau_{best}$ )
BM-25 kNN	0.036	0.043	0.718	0.805	0.069	0.134	0.113	0.002	0.001	0.001
LSTM	0.047	0.051	0.716	0.802	0.102	0.12	0.105	0.02	0.006	0.009
Transformer <sup>†</sup>	0.048									
T5 <sup>†</sup>	0.044									
CodeT5 <sup>†</sup>	0.048									
CodeReviewer	0.054	<b>0.071</b>	0.718	0.811	0.102	0.116	0.1	0.412	0.208	0.26
Stable-Code-Instruct-3B	0.042	0.04	0.733	0.784	0.091	0.16	0.133	0.34	0.199	0.228
Magicoder-S-DS-6.7B	0.035	0.041	0.72	0.815	0.101	0.175	0.151	0.445	0.491	0.42
DeepSeekCoder-Instruct-6.7B	0.045	0.054	<b>0.734</b>	0.782	<b>0.112</b>	0.183	0.157	0.513	0.418	0.422
CodeLLaMA-Instruct-7B	0.023	0.026	0.71	0.843	0.071	0.171	0.145	0.156	0.151	0.14
Llama-3-8B-Instruct	0.014	0.016	0.699	<b>0.898</b>	0.058	0.14	0.122	0.347	0.425	0.352
LLama-Reviewer <sup>†</sup>	<b>0.057</b>									
CodeLLaMA-Instruct-13B	0.025	0.029	0.715	0.839	0.079	0.179	0.152	0.274	0.272	0.253
GPT-3.5-Turbo	0.037	0.044	<b>0.734</b>	0.812	0.1	<b>0.2</b>	<b>0.171</b>	<b>0.563</b>	<b>0.635</b>	<b>0.558</b>

Table 12: Results of all eval. metrics and models on the entire test set. All metrics have been normalized to be between 0 and 1. <sup>†</sup> signifies reported scores. The reference based metrics have a very narrow range of values.

Smell Name	Description
Long method	There exist methods with too many lines (more lines than a set threshold).
Long parameter list	There exist methods with more than “n” parameters (n = 6 is used in this study).
Long branch	When conditional statement branches extend too long or are too nested.
Many attributes	When a single class has too many methods or attributes.
Many methods	When a single class has too many methods.
Shotgun surgery	When a single functionality is fragmented across various classes.
Class cohesion	There are some classes with low cohesion (C.2.1).
Code complexity	The code includes blocks with cyclomatic complexity (section C.2.2) of rank-C or worse (moderate to slightly complex blocks). Please read through the cyclomatic complexity and ranks section for more details.
Long lambda	The code includes lambda functions that exceed a threshold on length (number of characters).
Long list comprehension	The code includes list comprehensions that exceed a threshold on length (number of characters).

Table 13: Python code smells detected by the PyScent code smell static analysis tool

Rank	Human Annotated Relevance	Rel (ours)	chrF++ (best reference based metric)
1	GPT-3.5	GPT-3.5	GPT-3.5
2	LLaMA-3-8B-Instruct	DeepSeekCoder-6.7B-Instruct	DeepSeekCoder-6.7B-Instruct
3	DeepSeekCoder-6.7B-Instruct	Magocoder-S-DS-6.7B	CodeLLaMA-13B
4	Magocoder-S-DS-6.7B	LLaMA-3-8B-Instruct	Magocoder-S-DS-6.7B
5	CodeReviewer	CodeReviewer	Stable-Code-3B
6	Stable-Code-3B	Stable-Code-3B	LLaMA-3-8B-Instruct
7	CodeLLaMA-13B	CodeLLaMA-13B	BM-25 retriever
8	LSTM	LSTM	CodeReviewer
9	BM-25 retriever	BM-25 retriever	LSTM

Table 14: Rankings of the systems over the 300 human-annotated CodeReviewer instances according to the human-annotated relevance and our relevance metric. The rankings reveal that our method gets the rankings exactly right except for the LLaMA-3-8B-Instruct model (grayed out) which is ranked lower than Magocoder and DeepSeekCoder by our metric but preferred more by humans.

Language	Code Change	Claims	Review
Python	<pre> @@ -24,6 +24,7 @@ CREATE_TABLE = """     'complete_time' datetime DEFAULT NULL,     'status' enum('SUCCESS','RUNNING',     'FAILURE',     'PARTIAL_SUCCESS','TIMEOUT')     DEFAULT NULL,     + 'has_all_data' bool DEFAULT NULL,     'schema_version' varchar(255)     DEFAULT NULL,     'cycle_timestamp' varchar(255)     DEFAULT NULL,     PRIMARY KEY ('id'), </pre>	<p>A new column named 'has_all_data' has been added to the table. This column is of boolean data type and allows NULL values by default.</p> <p>The addition of the 'has_all_data' column implies that there might be a need to track whether all the data has been processed or not in the table. This could be used for further data processing or analysis.</p> <p>Existing rows in the table will not have a value for this new column. Therefore, it might require additional data updates or data entry processes to populate this column.</p>	<p><b>Magocoder:</b> This column is not being used anywhere in the codebase. It's a waste of space.</p>
Python	<pre> @@ -14,7 +14,9 @@ import abc import logging +import datetime import parameter +import target import warnings import traceback import parsing as pp </pre>	<p>Two new import statements have been added to the code. The first one imports the datetime module, and the second one imports the target module.</p> <p>The addition of the datetime module suggests that the code may now involve operations related to date and time. This could be for logging purposes, tracking the execution time of the code, or handling dates/times in the program.</p> <p>The addition of the target module indicates that the code may now involve operations related to the target environment or system. This could be for interacting with the target system, or for handling target-specific tasks.</p>	<p><b>CodeReviewer:</b> Why do we need these imports?</p>

Table 15: Cases where our metric underestimates the relevance of a review by scoring it as zero while the human scores it as 5 (max relevance). We observed that these cases tend to have fewer claims associated with the code change, briefer reviews with few relevant tokens, and fewer inline code snippets.

Language	Code Change	Claims	Review
Python	<pre> @@ -971,7 +971,7 @@ class JobTask(BaseHadoopJobTask):     if self.__module__ == '__main__':         d = pickle.dumps(self)         module_name = os.path.basename(             sys.argv[0]).rsplit('.', 1)[0]         d = d.replace(b'(__main__', "(c" + module_name) +       d = d.replace(b'__main__', b'c' + \         module_name.encode('ascii'))         open(file_name, "wb").write(d)      else: </pre>	<p>The code change is in the replacement of a part of the byte string 'c__main__' with 'c' followed by the encoded ASCII representation of 'module_name'.</p> <p>The original code was replacing 'c__main__' with 'c' followed by 'module_name'.</p> <p>The new code change corrects this by encoding 'module_name' to ASCII before replacing 'c__main__' with 'c' followed by the encoded 'module_name'.</p> <p>This change could potentially affect the behavior of the code if 'module_name' contains non-ASCII characters.</p>	<p><b>CodeReviewer:</b> Shouldn't this be 'module_name.encode('ascii')'?</p> <p><b>StableCode:</b> This seems unnecessary, the pickle module already encodes the module name.</p>

Table 16: Cases where our metric overestimates the relevance of a review by assigning it the maximum score while the human scores it at the lowest Likert score value of 1. We observed that these cases tend to have far more inline-code snippets than the rest of the data and underestimation cases.